

CS 2051: Honors Discrete Mathematics

Spring 2023 Homework 5 Supplement

Sarthak Mohanty

Overview

Note: This document and the autograder will be continuously updated until Monday morning, but the fundamentals remain the same.

In this supplement, you learn about a generalized form of functions known as **relations**, and explore the connection between functions and relations. You'll learn about equivalence relations and partial orders. You'll also explore some of the applications of these concepts

Part 1: Generalizing Functions with Relations

This week, you learned about functions. We (informally) defined them as a well-defined mapping between two sets. However, what about mappings that are not well-defined? Is there no way to represent these? As a matter of fact, there is!

Definition a *relation* \mathcal{R} over sets A, B is a subset of $A \times B$. The notation $a\mathcal{R}b$ or $a \sim b$ is often used to denote that $(a, b) \in \mathcal{R}$.

Examples

Let $\mathcal{R}_1, \dots, \mathcal{R}_4$ be a relation on $A = \{1, 2, 3, 4\}$.

- $\mathcal{R}_1 = \{(a, b) \mid a \leq b\}$
- $\mathcal{R}_2 = \{(a, b) \mid a = b\}$
- $\mathcal{R}_3 = \{(a, b) \mid a + b \leq 2022\}$
- $\mathcal{R}_4 = \{(a, b) \mid a \text{ divides } b\}$

Properties of a Relation

Reflexivity \mathcal{R} is *reflexive* if

$$(\forall a \in A)(a\mathcal{R}a)$$

Symmetry \mathcal{R} is *symmetric* if

$$(\forall a, b \in A)(a\mathcal{R}b \iff b\mathcal{R}a)$$

Antisymmetry \mathcal{R} is *antisymmetric* if

$$(\forall a, b \in A)(a\mathcal{R}b \wedge b\mathcal{R}a \rightarrow a = b)$$

Transitivity \mathcal{R} is *transitive* if

$$(\forall a, b, c \in A)(a\mathcal{R}b \wedge b\mathcal{R}c \rightarrow a\mathcal{R}c)$$

When a relation is reflexive, antisymmetric, and transitive, we call it a *partial order*.

When a relation is reflexive, symmetric, and transitive, we call it a *equivalence relation*.

In this part, you'll implement the functions `isPartialOrder(elements, relation)` and `isEquivalenceRelation(elements, relation)`. These functions takes in a relation (represented as a list of tuples) and returns whether or not the relation (taken over the set of elements) is a valid partial order or equivalence relation, respectively. You must use the following helper methods:

- `isReflexive(elements, relation)`
- `isSymmetric(elements, relation)`
- `isAntisymmetric(elements, relation)`
- `isTransitive(elements, relation)`

All methods (and all helper methods) must be implemented in one line. Hint: use `all` method

Part 2: Partitioning with Equivalence Relations

In the world of computer science, there are two main applications for relations: partitioning and scheduling. In this part, we'll cover partitioning, which is essentially just a reframing of our knowledge about equivalence relations. If any of the concepts introduced in this section feel rather hand-wavy, feel free to consult the textbook, which has rigorous proofs for the theorems.

Definition: Given some relation \mathcal{R} over the set A , the *equivalence class* of an element $x \in A$ is

$$[x] = \{y : x\mathcal{R}y\}$$

There is a very powerful theorem related to this concept:

Important Theorem: The equivalence classes of an equivalence relation on a set A *partition* A into a collection of disjoint, nonempty subsets A_1, A_2, \dots, A_n such that (and this is the important part) $\bigcup_{i=1}^n A_i = A$.

Ok, so this fact is cool and all, but what's the point? Well, if we have say

Example 1: Congruence Relations

Our first example delves into *number theory*, a field you will become more intimate with in the next few supplements. Informally, define the relation $a\mathcal{R}_n b$ over $\mathbb{Z} \times \mathbb{Z}$ if a and b have the same remainder when divided by some number n . We usually denote this using

$$a \equiv b \pmod{n}.$$

For example, -10 and 15 are related under \mathcal{R}_5 ,

$$-10 \equiv 15 \pmod{5}$$

since $-10 - 15 = -25$ is a multiple of 5 , or equivalently since both -10 and 15 have the same remainder 0 when divided by 5 .

All such relations \mathcal{R}_n are equivalence relations, and partition the set of integers into n equivalence classes. For example, the relation \mathcal{R}_3 partitions the integers like so

$$\{\dots, -8, -5, -2, 1, 4, 7, \dots\} \tag{1}$$

$$\{\dots, -7, -4, -1, 2, 5, 8, \dots\} \tag{2}$$

$$\{\dots, -6, -3, 0, 3, 6, 9, \dots\} \tag{3}$$

The canonical

Example 2: Vector Spaces

One of the most prominent uses of equivalence classes is in linear algebra. After all, it's very difficult to determine when two vector spaces or matrices are functionally “identical”.

If you've taken an introductory linear algebra course, you learned about Gauss's Method to solve a system of equations. In essence, it worked by starting with a matrix and deriving a sequence of other matrices, each *row equivalent* to each other. It turns out that row equivalence is an equivalence relation, and partitions the set of all matrices into corresponding classes, as shown in Figure (2).

We can generalize this one step further with a relation known as *matrix equivalence*. Matrix equivalent matrices represent the same map, with respect to appropriate pairs of bases. Matrix equivalence classes are characterized by rank: two same-sized matrices are matrix equivalent if and only if they have the same rank. In fact, matrix similarity is a special case of matrix equivalence!

The canonical forms for row equivalence are the Reduced Echelon form matrices, which you may already be familiar with. Meanwhile, the canonical forms for matrix similarity are Jordan Normal forms. (the canonical forms for matrix equivalence are a bit more complicated.)





Figure 1: Matrix similarity is a finer partition than matrix equivalence



(a) Two matrices with the same row space.

(b) The canonical forms for the set of 2×2 matrices.

Figure 2: Row equivalence as an equivalence relation.

In this part, you'll implement the following function:

- `partition(elements, relation)`: This function takes an equivalence relation (this time represented as a boolean function), and returns a partition of the elements into equivalence classes. For example, given the congruence relation described above, the function should return

```
>>> partition([i for i in range(-8, 8)], lambda x, y: (x - y) % 3 == 0)
[{-6, -3, 0, 3, 6}, {-8, -5, -2, 1, 4, 7}, {-7, -4, -1, 2, 5, 8}]
```

Tip: iterating over sets is difficult. Try casting.

Part 3: Single Processor Task Scheduling


We now delve into partial orderings, or “posets”. We’ll cover a nice “graph”ical representation of posets, and then cover their applications in the context of task scheduling.

Dependency Graphs

To start, let’s consider the following story.

You’re a recently admitted CS major at Georgia Tech. To save money, you’re trying to graduate as early as possible. You have a lot of course that you can take, but you’re not sure which ones to take when. Furthermore, many of the courses are dependent upon completion of other courses (for example, CS 2110 requires CS 1331, which in turn requires CS 1301). How long will it take for you to graduate?

We can represent this problem as a poset \mathcal{R} over the set of courses A , where $a\mathcal{R}b$ iff a is a prerequisite to take b :



This is known as a *dependency graph*.

Task Scheduling

Using the formalism of posets, we can also define the problem of computing an optimal schedule. We do so in the language of job scheduling, as it’s the most common form.

Assume that some task T can be decomposed into sub-tasks $T = \{T_1, T_2, \dots, T_n\}$. In general, we can encode the relationships between the various T_i as a poset, where $T_i < T_j$ if (sub)task T_i needs to be performed before T_j . **Assume each task T_i takes 1 unit of time to complete.** Given k processors P_1, P_2, \dots, P_k

a **schedule** for T assigns each sub-task T_i a processor P_j as well as a time t_i at which P_j should start task T_i . Observe that if a processor starts T_i at times t_i , then the task will complete at time $t_i + w_i$, where w_i is the weight of T_i . A schedule is **feasible** if:

- 1 no single processor is performing multiple tasks at the same time (i.e., if a processor is assigned T_i at time t_i and T_j at time $t_j > t_i$, then $t_j \geq t_i + w_i$), and
- 2 for every pair of tasks T_i and T_j with $T_i \prec T_j$, T_j is schedule to start some time after (or at the same time) T_i completes (i.e., $t_j \geq t_i + w_i$).

If we have only one processor, then the answer is fairly simple, just create a topological sorting: Definition: A topological sorting is a total ordering of a partially ordered set.

Algorithm 1 TOPOLOGICAL_SORT($poset$)

```

 $T$  = empty list
 $S$  = all minimal elements in  $poset$ 
while  $S$  is not empty do
    remove some node  $u$  from  $S$  and add it to  $T$ .
     $uv\_dependencies$  set of all dependencies of the form  $(u, v)$  for some  $v$ .
    for each  $(u, v)$  in  $uv\_dependencies$  do
        if  $v$  is minimal then Insert  $v$  in  $T$ 
return  $T$ 

```

In this part, you'll implement the following function:

`topological_sort(poset)`: This method takes in a partially ordered set in the form of a dependency list and returns a valid topological sort.

I'll be manually checking that you used some semblance of Kahn's algorithm in your solution, so don't use other approaches like a modified DFS (besides, it'll be more useful for the next part).

Tip: The use of external data structures such as queues or stacks may be helpful (but not required) in completing this task.

Part 4: Multi-Processor Job Scheduling

Definition A *chain* is such that every distinct pair of elements is comparable in \mathcal{P} .

Note that the time it takes to schedule tasks, even with an unlimited number of processors, is at least the length of the longest chain. Indeed, if we used less time, then two items from a longest chain would have to be done at the same time, which contradicts the precedence constraints.

The nice thing about posets is that this is always possible! In other words, for any poset, there is a legal parallel schedule that runs in t steps, where t is the length of the longest chain. This is known as Mirsky's theorem:

If we have an unlimited number of processors, then the time to complete all tasks is equal to the length of the longest chain of dependent tasks. The case where there are a limited number of processors, however, is more useful in practice.

`generate_schedule(poset, num_processors)`: This method takes in a partially ordered set in the form of a dependency list and returns a valid schedule in the form of a list of lists, where the i -th element in the list represents the jobs we should schedule at time $t = i$.

Submission Instructions (10 pts)

After you fill the appropriate functions, submit the following files to Gradescope and make sure you pass all test cases:

- `relations.py`
- `scheduler.py`

Notes

- The autograder may not reflect your final grade on the assignment. We reserve the right to run additional tests during grading.
- Do not import additional packages, as your submission may not pass the test cases or manual review.